

# co2-emissions-simple

October 7, 2023

## 1 Act 1. Regresion lineal Simple

CO2 Emission by Vehicles

- Andrés Alejandro Guzmán González | A01633819
- Joel Isaias Solano Ocampo | A01639289
- Tania Sayuri Guizado Hernández | A01640092
- Ernesto Reynoso Lizárraga | A01639915

---

### Contenidos

- 5.1 - Temas avanzados sobre análisis regresión
- 5.1.1 - Verificación de supuestos: QQ-plots & análisis de residuales
- 5.2 - Estadística para datos multivariados

### 1. ¿Cuáles son las características que más influyen en la emisión de CO2?

Considerando las pruebas realizadas en esta actividad, las variables con más influencias en la emisión de Dioxido de Carbono (CO2) son:

- \* Fuel Consumption Comb (mpg)
- \* Fuel Consumption Comb (L/100 km)
- \* Fuel Consumption City (L/100 km)

Pues su correlación es mayor a 0.8

### 2. ¿Habrá alguna diferencia en la emisiones de CO2 cuando el consumo de combustible para la ciudad y carretera se consideran por separado?

Si, ya que el consumo de combustible en la ciudad tiene una mayor correlación con el gasto total. Al combinarlos es más cercana la predicción a las emisiones de Dioxido de Carbono (CO2) totales. Esto se da ya que al suponer que las emisiones Pues así al combiarlas el consumo

##Llamado a librerías

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from scipy.stats import norm, uniform, skewnorm
```

##Importamos los datos a un DataFrame

```
[ ]: df = pd.read_csv('/content/drive/Shared drives/Reto IA/Actividades/CO2_Emissions/
↳CO2_Emissions_Canada.csv')
df.head()
```

```
[ ]:      Make      Model Vehicle Class  Engine Size(L)  Cylinders Transmission \
0  ACURA      ILX      COMPACT          2.0           4           AS5
1  ACURA      ILX      COMPACT          2.4           4           M6
2  ACURA  ILX HYBRID      COMPACT          1.5           4           AV7
3  ACURA      MDX 4WD    SUV - SMALL          3.5           6           AS6
4  ACURA      RDX AWD    SUV - SMALL          3.5           6           AS6
```

```
      Fuel Type  Fuel Consumption City (L/100 km) \
0             Z              9.9
1             Z             11.2
2             Z              6.0
3             Z             12.7
4             Z             12.1
```

```
      Fuel Consumption Hwy (L/100 km)  Fuel Consumption Comb (L/100 km) \
0              6.7              8.5
1              7.7              9.6
2              5.8              5.9
3              9.1             11.1
4              8.7             10.6
```

```
      Fuel Consumption Comb (mpg)  CO2 Emissions(g/km)
0              33              196
1              29              221
2              48              136
3              25              255
4              27              244
```

```
[ ]: df.isnull().sum()
```

```
[ ]: Make
      Model
      Vehicle Class
      Engine Size(L)
      Cylinders
      Transmission
      Fuel Type
      Fuel Consumption City (L/100 km)
      Fuel Consumption Hwy (L/100 km)
      Fuel Consumption Comb (L/100 km)
      Fuel Consumption Comb (mpg)
      CO2 Emissions(g/km)
      dtype: int64
```

```
[ ]: df.shape
```

```
[ ]: (7385, 12)
```

```
[ ]: df.columns
```

```
[ ]: Index(['Make', 'Model', 'Vehicle Class', 'Engine Size(L)', 'Cylinders',
           'Transmission', 'Fuel Type', 'Fuel Consumption City (L/100 km)',
           'Fuel Consumption Hwy (L/100 km)', 'Fuel Consumption Comb (L/100 km)',
           'Fuel Consumption Comb (mpg)', 'CO2 Emissions(g/km)'],
          dtype='object')
```

```
[ ]: df.describe()
```

```
[ ]:      Engine Size(L)      Cylinders  Fuel Consumption City (L/100 km) \
count      7385.000000    7385.000000                7385.000000
mean         3.160068         5.615030                12.556534
std          1.354170         1.828307                 3.500274
min           0.900000         3.000000                 4.200000
25%           2.000000         4.000000                10.100000
50%           3.000000         6.000000                12.100000
75%           3.700000         6.000000                14.600000
max           8.400000        16.000000                30.600000

      Fuel Consumption Hwy (L/100 km)  Fuel Consumption Comb (L/100 km) \
count                7385.000000                7385.000000
mean                  9.041706                 10.975071
std                   2.224456                 2.892506
min                   4.000000                 4.100000
25%                   7.500000                 8.900000
50%                   8.700000                10.600000
75%                  10.200000                12.600000
max                  20.600000                26.100000
```

	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
count	7385.000000	7385.000000
mean	27.481652	250.584699
std	7.231879	58.512679
min	11.000000	96.000000
25%	22.000000	208.000000
50%	27.000000	246.000000
75%	32.000000	288.000000
max	69.000000	522.000000

##Funciones genereales

```
[ ]: def plt_scatter(x, y, xl, yl):
    plt.subplot(1, 2, 1)
    plt.scatter(x, y)
    plt.xlabel(xl)
    plt.ylabel(yl)
    plt.grid()
```

```
[ ]: def plt_scatter_w_line(x, y, b0, b1, xl, yl):
    x_line = np.linspace(min(x), max(x), 100)
    y_line = b0 + b1 * x_line
    plt.subplot(1, 2, 2)
    plt.scatter(x,y)
    plt.xlabel(xl)
    plt.ylabel(yl)
    plt.grid()

    plt.plot(x_line, y_line, color='red')
```

```
[ ]: def Asimetric_Dist(r, title):
    plt.hist(r, density = True, bins = 'auto', histtype='stepfilled', alpha=0.2)
    plt.title(title)
    plt.grid()
    plt.show()
```

```
[ ]: def OLS(x, y):
    x = sm.add_constant(x)
    model = sm.OLS(y, x)
    result = model.fit()
    print('Params:', result.params)
    print('R^2:', result.rsquared)
```

## 1.1 Engine Size(L)

```
[ ]: x_ms = df['Engine Size(L)']  
y = df['CO2 Emissions(g/km)']
```

```
[ ]: X_ms = sm.add_constant(x_ms)  
print(X_ms.shape)  
print(X_ms)
```

```
(7385, 2)  
      const  Engine Size(L)  
0      1.0      2.0  
1      1.0      2.4  
2      1.0      1.5  
3      1.0      3.5  
4      1.0      3.5  
...    ...  
7380    1.0      2.0  
7381    1.0      2.0  
7382    1.0      2.0  
7383    1.0      2.0  
7384    1.0      2.0
```

```
[7385 rows x 2 columns]
```


```
[ ]: model = sm.OLS(y,X_ms)  
result = model.fit()  
result.params
```

```
[ ]: const      134.365893  
      Engine Size(L)    36.777315  
      dtype: float64
```

```
[ ]: print("\nR2: ", result.rsquared)
```

```
R2:  0.7244472046524082
```

```
[ ]: plt.figure(figsize=(12, 6))  
plt.scatter(x_ms, y, 'Engine Size(L)', 'CO2 Emissions(g/km)')  
plt.scatter_w_line(x_ms, y, result.params[0], result.params[1], 'Engine_  
↪Size(L)', 'CO2 Emissions(g/km)')  
plt.tight_layout()  
plt.show()
```



## 1.2 Cylinders


```
[ ]: x_cyl = df['Cylinders']
      X_cyl = sm.add_constant(x_cyl)
      model = sm.OLS(y,X_cyl)
      result = model.fit()
      result.params
```

```
[ ]: const      100.956915
      Cylinders   26.647724
      dtype: float64
```

```
[ ]: print("\nR2: ", result.rsquared)
```

```
R2:  0.6932953649936133
```

```
[ ]: plt.figure(figsize=(12, 6))
      plt.scatter(x_cyl, y, 'Cylinders', 'CO2 Emissions(g/km)')
      plt.scatter_w_line(x_cyl, y, result.params[0], result.params[1], 'Cylinders',
                          ↪ 'CO2 Emissions(g/km)')
      plt.tight_layout()
      plt.show()
```




### 1.3 Fuel Consumption Hwy (L/100 km)

```
[ ]: x_fch = df['Fuel Consumption Hwy (L/100 km)']
      X_fch = sm.add_constant(x_fch)
      model = sm.OLS(y,X_fch)
      result = model.fit()
      print(result.params)
      print("\nR2: ", result.rsquared)
```

```
const                40.448581
Fuel Consumption Hwy (L/100 km)  23.240759
dtype: float64
```


```
R2:  0.7806357669286315
```

```
[ ]: plt.figure(figsize=(12, 6))
      plt.scatter(x_fch, y, 'Fuel Consumption Hwy (L/100 km)', 'CO2 Emissions(g/km)')
      plt.scatter_w_line(x_fch, y, result.params[0], result.params[1], 'Fuel_
      ↪Consumption Hwy (L/100 km)', 'CO2 Emissions(g/km)')
      plt.tight_layout()
      plt.show()
```




##Análisis de distribuciones

```
[ ]: y_skew = skewnorm.rvs(y)
      Asimetric_Dist(y_skew, 'Original - Positive Skew')
```







```
[ ]: y_root = np.sqrt(y + abs(min(y)))
      Asimetric_Dist(y_root, 'Root Squared - Transformed Data')
```




```
[ ]: y_log = np.log10(1 + y + abs(min(y)))
      Asimetric_Dist(y_log, 'Log - Transformed Data')
```




```
[ ]: y_skew_neg = -skewnorm.rvs(y)
      Asimetric_Dist(y_skew_neg, 'Original - Negative Skew')
```



```
[ ]: y_pos = y + abs(min(y))  
     y_root_neg = np.sqrt(max(y_pos) - y_pos)  
     Asimetric_Dist(y_root_neg, 'Root Squared - Transformed Data')
```



```
[ ]: y_positive = y + abs(min(y))  
      y_log_neg = np.log10(1 + max(y_positive) - y_positive)  
      Asimetric_Dist(y_log_neg, 'Log - Transformed Data')
```



#### 1.4 Fuel Consumption City (L/100 km)


```
[ ]: x_fcct = df['Fuel Consumption City (L/100 km)']
      X_fcct = sm.add_constant(x_fcct)
      model = sm.OLS(y,X_fcct)
      result = model.fit()
      print(result.params)
      print("\nR2: ", result.rsquared)
```

```
const                    57.559903
Fuel Consumption City (L/100 km)  15.372459
dtype: float64
```

```
R2:  0.8456503198972763
```

```
[ ]: plt.figure(figsize=(12, 6))
      plt.scatter(x_fcct, y, 'Fuel Consumption City (L/100 km)', 'CO2 Emissions(g/
      ↪km)')
      plt.scatter_w_line(x_fcct, y, result.params[0], result.params[1], 'Fuel_
      ↪Consumption City (L/100 km)', 'CO2 Emissions(g/km)')
```

```
plt.tight_layout()
plt.show()
```



```
[ ]: print(result.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:      CO2 Emissions(g/km)    R-squared:                0.846
Model:              OLS                    Adj. R-squared:           0.846
Method:             Least Squares          F-statistic:             4.045e+04
Date:               Sat, 07 Oct 2023        Prob (F-statistic):      0.00
Time:               02:13:24                Log-Likelihood:          -33630.
No. Observations:   7385                    AIC:                     6.726e+04
Df Residuals:       7383                    BIC:                     6.728e+04
Df Model:           1
Covariance Type:    nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                57.5599      0.996    57.772    0.000
55.607    59.513
Fuel Consumption City (L/100 km)  15.3725      0.076   201.122    0.000
15.223    15.522
=====
```

```
=====
Omnibus:             3089.403    Durbin-Watson:           1.913
Prob(Omnibus):       0.000     Jarque-Bera (JB):        16424.392
```

Skew:	-1.963	Prob(JB):	0.00
Kurtosis:	9.161	Cond. No.	48.8

=====


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[ ]: x_line = np.linspace(min(x_fcct), max(x_fcct), 100)
y_line = result.params[0] + result.params[1] * x_line
plt.scatter(x_fcct,y)
plt.xlabel('Fuel Consumption City (L/100 km)')
plt.ylabel('CO2 Emissions (g/km)')
plt.grid()

plt.plot(x_line, y_line, color='red')

plt.show()
```




```
[ ]: influence = result.get_influence()
std_residual = influence.resid_studentized_internal
```

```
print(std_residual)
```


```
[-0.5980394 -0.37982853 -0.60022109 ...  0.11233374  0.09868503  
 0.12598264]
```

```
[ ]: plt.scatter(x_fcct, std_residual)  
plt.xlabel('Fuel Consumption City (L/100 km)')  
plt.ylabel('standardized residual')  
plt.axhline(y=0, color='red', linestyle='--', linewidth=1)  
plt.show()
```




```
[ ]: fig = sm.qqplot(std_residual, dist=norm, line="q")  
plt.ylabel("standardized residuals Quantiles")  
plt.title("QQ-Plot - normal distribution")  
plt.grid()
```

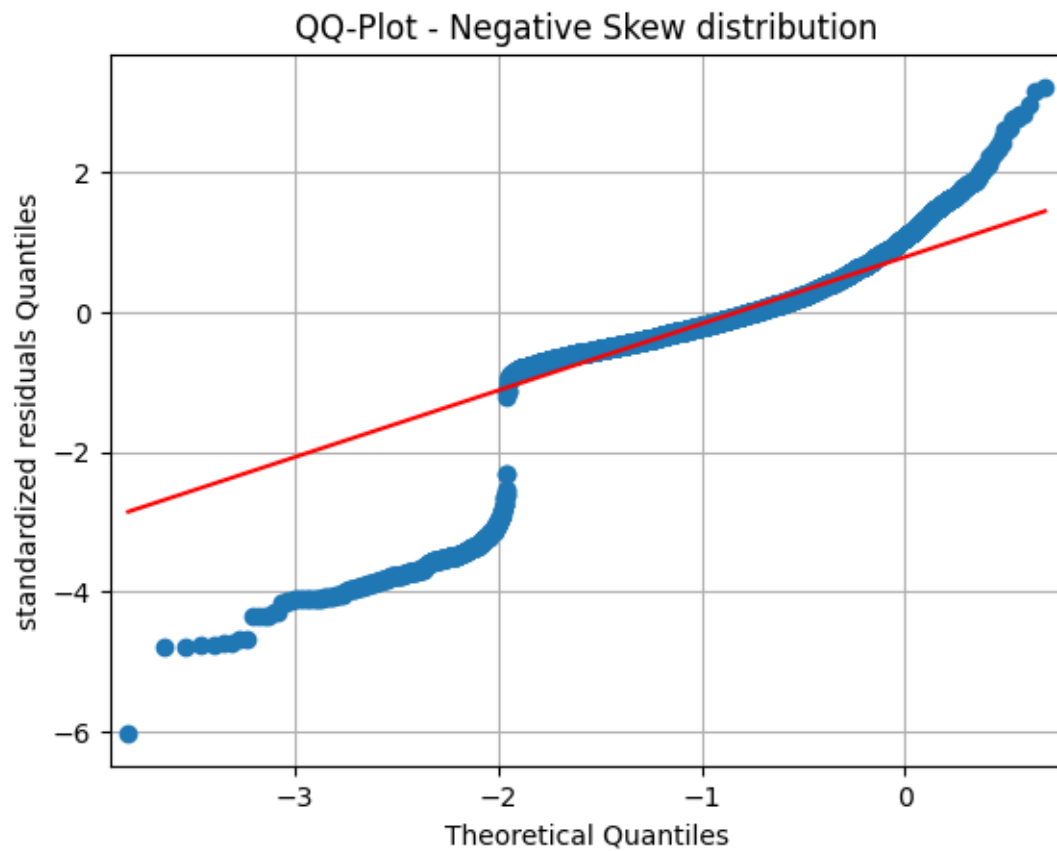




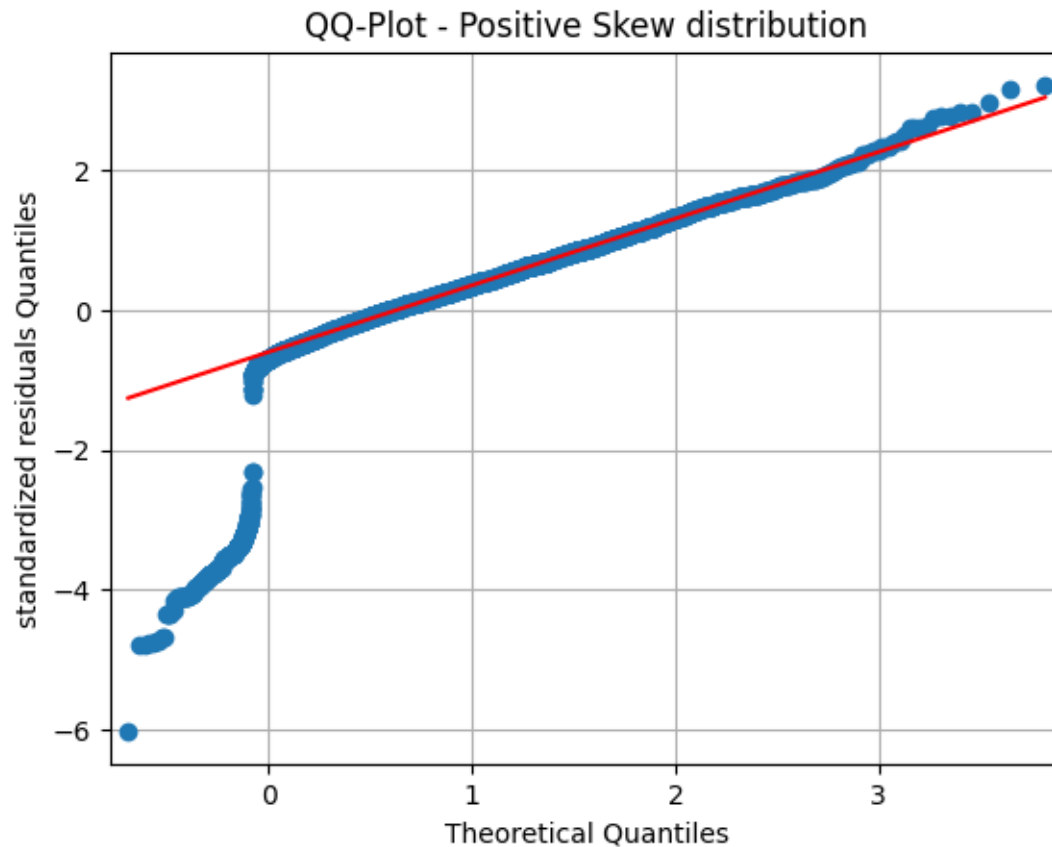
```
[ ]: fig = sm.qqplot(std_residual, dist=uniform, line="q")  
plt.ylabel("standardized residuals Quantiles")  
plt.title("QQ-Plot - uniform distribution")  
plt.grid()
```



```
[ ]: fig = sm.qqplot(std_residual, skewnorm(-4), line="q")
plt.ylabel("standardized residuals Quantiles")
plt.title("QQ-Plot - Negative Skew distribution")
plt.grid()
```



```
[ ]: fig = sm.qqplot(std_residual, skewnorm(4), line="q")
plt.ylabel("standardized residuals Quantiles")
plt.title("QQ-Plot - Positive Skew distribution")
plt.grid()
```



```
[ ]: OLS(x_fcct, y_root)
```

```
Params: const                                13.419645
Fuel Consumption City (L/100 km)            0.408706
dtype: float64
R^2: 0.844322646434444
```

```
[ ]: OLS(x_fcct, y_log)
```

```
Params: const                                2.296990
Fuel Consumption City (L/100 km)            0.018955
dtype: float64
R^2: 0.836968070475829
```

## 1.5 Fuel Consumption Comb (mpg)


```
[ ]: x_fcc = df['Fuel Consumption Comb (mpg)']
X_fcc = sm.add_constant(x_fcc)
model = sm.OLS(y,X_fcc)
result = model.fit()
```

```
print(result.params)
print("\nR2: ", result.rsquared)
```

```
const                452.353036
Fuel Consumption Comb (mpg)    -7.341929
dtype: float64
```

```
R2: 0.8234224657110062
```

```
[ ]: plt.figure(figsize=(12, 6))
plt.scatter(x_fcc, y, 'Fuel Consumption Comb (mpg)', 'CO2 Emissions(g/km)')
plt.scatter_w_line(x_fcc, y, result.params[0], result.params[1], 'Fuel_
Consumption Comb (mpg)', 'CO2 Emissions(g/km)')
plt.tight_layout()
plt.show()
```



```
[ ]: print(result.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:    CO2 Emissions(g/km)    R-squared:                0.823
Model:                OLS                Adj. R-squared:           0.823
Method:                Least Squares      F-statistic:             3.443e+04
Date:                Sat, 07 Oct 2023     Prob (F-statistic):       0.00
Time:                02:13:29             Log-Likelihood:          -34127.
No. Observations:    7385                AIC:                    6.826e+04
Df Residuals:        7383                BIC:                    6.827e+04
Df Model:              1
Covariance Type:      nonrobust
```

```
=====
=====
                                coef      std err          t      P>|t|
[0.025      0.975]
-----
const                452.3530      1.124      402.297      0.000
450.149      454.557
Fuel Consumption Comb (mpg)  -7.3419      0.040     -185.550      0.000
-7.419      -7.264
=====
Omnibus:                1935.010      Durbin-Watson:                1.326
Prob(Omnibus):            0.000      Jarque-Bera (JB):            13170.162
Skew:                    1.080      Prob(JB):                    0.00
Kurtosis:                9.176      Cond. No.                    112.
=====
```

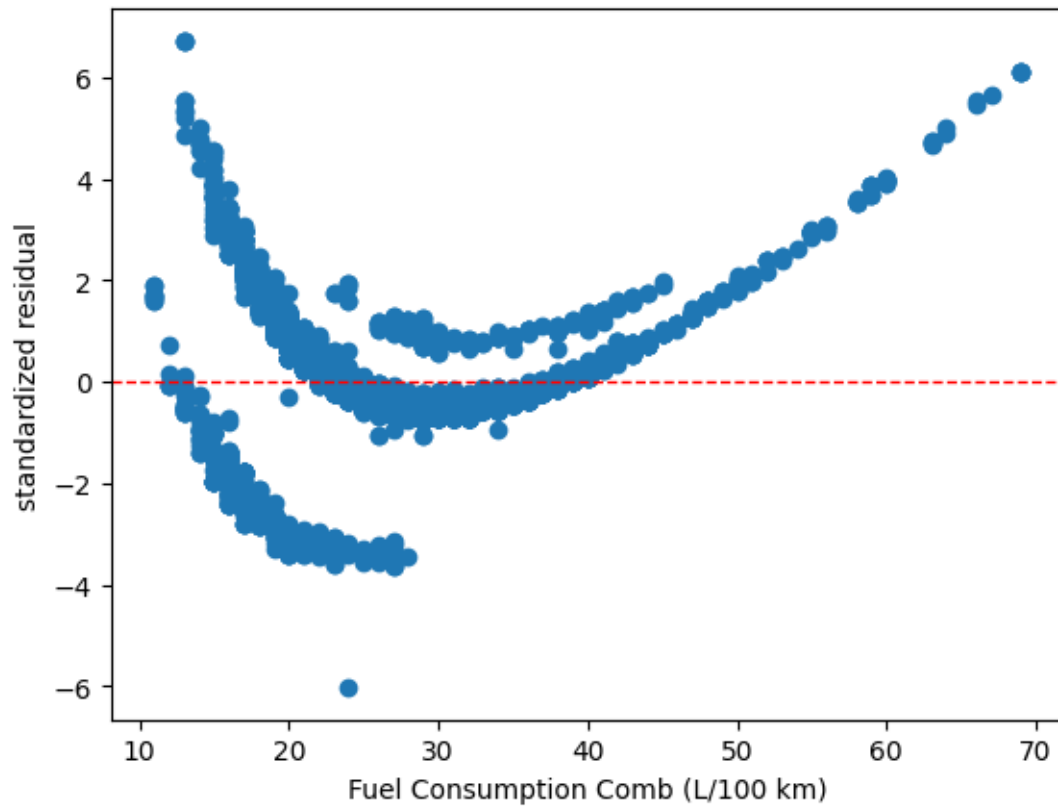
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```
[ ]: influence = result.get_influence()
std_residual = influence.resid_studentized_internal
print(std_residual)
```

```
[-0.57223492 -0.74985292  1.46736968 ... -0.57431008 -0.30247323
-0.54754722]
```

```
[ ]: plt.scatter(x_fcc, std_residual)
plt.xlabel('Fuel Consumption Comb (L/100 km)')
plt.ylabel('standardized residual')
plt.axhline(y=0, color='red', linestyle='--', linewidth=1)
plt.show()
```




```
[ ]: fig = sm.qqplot(std_residual, dist=norm, line="q")  
plt.ylabel("standardized residuals Quantiles")  
plt.title("QQ-Plot - Normal distribution")  
plt.grid()
```




```
[ ]: fig = sm.qqplot(std_residual, dist=uniform, line="q")  
plt.ylabel("standardized residuals Quantiles")  
plt.title("QQ-Plot - Uniform distribution")  
plt.grid()
```






```
[ ]: fig = sm.qqplot(std_residual, skewnorm(-4), line="q")
plt.ylabel("standardized residuals Quantiles")
plt.title("QQ-Plot - Negative Skew distribution")
plt.grid()
```



```
[ ]: fig = sm.qqplot(std_residual, skewnorm(4), line="q")  
plt.ylabel("standardized residuals Quantiles")  
plt.title("QQ-Plot - Positive Skew distribution")  
plt.grid()
```



```
[ ]: OLS(x_fcc, y_root)
```

```
Params: const                24.024339
Fuel Consumption Comb (mpg)  -0.199142
dtype: float64
R^2: 0.8556794221893714
```

```
[ ]: OLS(x_fcc, y_log)
```

```
Params: const                2.793990
Fuel Consumption Comb (mpg)  -0.009424
dtype: float64
R^2: 0.8831148871092463
```

## 1.6 Fuel Consumption Comb (L/100 km)


```
[ ]: x_fccl = df['Fuel Consumption Comb (L/100 km)']
X_fccl = sm.add_constant(x_fccl)
model = sm.OLS(y,X_fccl)
result = model.fit()
```

```
print(result.params)
print("\nR2: ", result.rsquared)
```

```
const                                46.763152
Fuel Consumption Comb (L/100 km)     18.571319
dtype: float64
```

```
R2: 0.8428186895623988
```

```
[ ]: plt.figure(figsize=(12, 6))
plt.scatter(x_fccl, y, 'Fuel Consumption Comb (L/100 km)', 'CO2 Emissions(g/
↪km)')
plt.scatter_w_line(x_fccl, y, result.params[0], result.params[1], 'Fuel_
↪Consumption Comb (L/100 km)', 'CO2 Emissions(g/km)')
plt.tight_layout()
plt.show()
```



```
[ ]: print(result.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:    CO2 Emissions(g/km)    R-squared:                0.843
Model:            OLS                    Adj. R-squared:           0.843
Method:           Least Squares          F-statistic:             3.959e+04
Date:             Sat, 07 Oct 2023        Prob (F-statistic):       0.00
Time:             02:13:27                Log-Likelihood:          -33697.
No. Observations: 7385                    AIC:                     6.740e+04
Df Residuals:     7383                    BIC:                     6.741e+04
Df Model:         1
```

```

Covariance Type:            nonrobust
=====
=====
                                coef      std err          t      P>|t|
[0.025      0.975]
-----
const                        46.7632      1.059      44.142      0.000
44.686      48.840
Fuel Consumption Comb (L/100 km)  18.5713      0.093     198.968      0.000
18.388      18.754
=====
Omnibus:                    3592.018   Durbin-Watson:           1.986
Prob(Omnibus):              0.000   Jarque-Bera (JB):       22309.895
Skew:                      -2.290   Prob(JB):               0.00
Kurtosis:                  10.178   Cond. No.               44.9
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

[ ]: influence = result.get_influence()
std_residual = influence.resid_studentized_internal
print(std_residual)

```

```


[-0.37157485 -0.17449253 -0.87672126 ...  0.0841568  0.05952249
 0.10879112]

```


```

[ ]: plt.scatter(df['Fuel Consumption Comb (L/100 km)'], std_residual)
plt.xlabel('Fuel Consumption Comb (L/100 km)')
plt.ylabel('standardized residual')
plt.axhline(y=0, color='red', linestyle='--', linewidth=1)
plt.show()

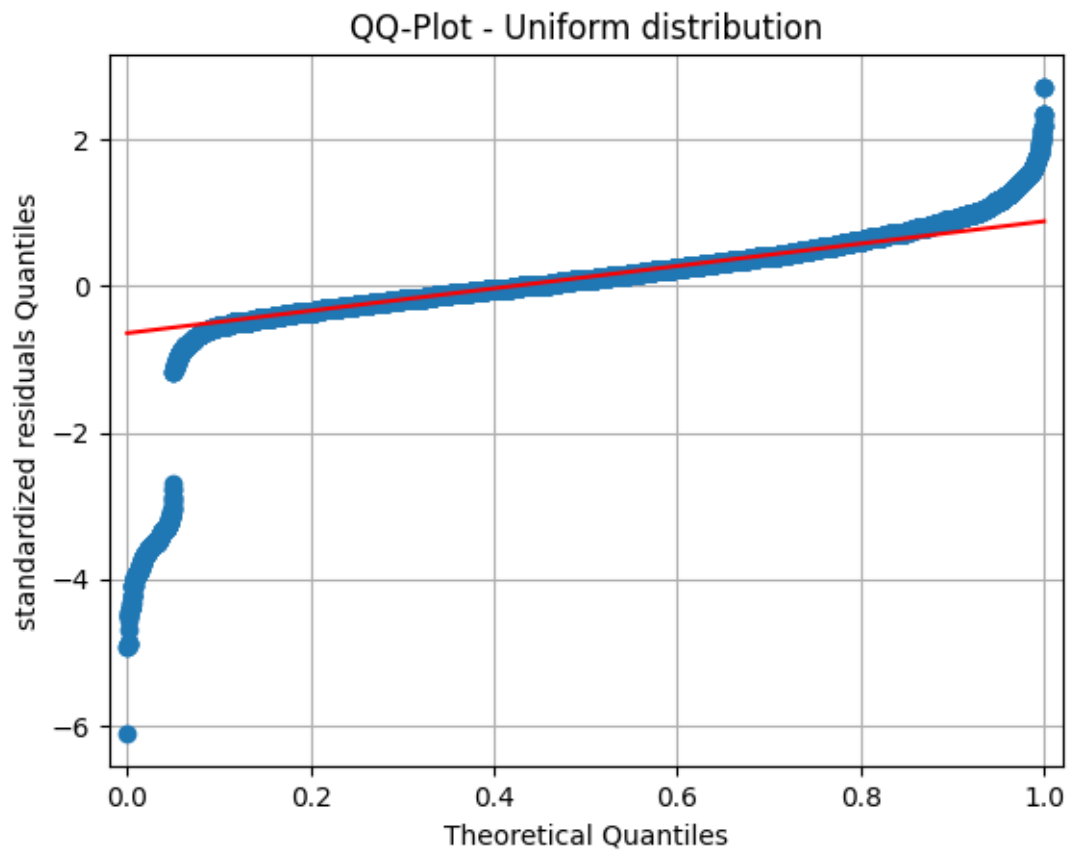
```



```
[ ]: fig = sm.qqplot(std_residual, dist=norm, line="q")
plt.ylabel("standardized residuals Quantiles")
plt.title("QQ-Plot - Normal distribution")
plt.grid()
```



```
[ ]: fig = sm.qqplot(std_residual, dist=uniform, line="q")  
plt.ylabel("standardized residuals Quantiles")  
plt.title("QQ-Plot - Uniform distribution")  
plt.grid()
```




```
[ ]: fig = sm.qqplot(std_residual, skewnorm(-4), line="q")
plt.ylabel("standardized residuals Quantiles")
plt.title("QQ-Plot - Negative Skew distribution")
plt.grid()
```





```
[ ]: fig = sm.qqplot(std_residual, skewnorm(4), line="q")  
plt.ylabel("standardized residuals Quantiles")  
plt.title("QQ-Plot - Positive Skew distribution")  
plt.grid()
```



```
[ ]: OLS(x_fccl, y_root)
```

```
Params: const                                13.134986
Fuel Consumption Comb (L/100 km)            0.493536
dtype: float64
R^2: 0.8407523790886398
```

```
[ ]: OLS(x_fccl, y_log)
```

```
Params: const                                2.283923
Fuel Consumption Comb (L/100 km)            0.022877
dtype: float64
R^2: 0.8325284987248827
```